

EN3160 Assignment 1 on Intensity Transformations and Neighborhood Filtering

Github Link: https://github.com/oshan-imaduwage/EN3160_Assignment_01


```
In [17]: import cv2
import numpy as np
import math
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = (5, 3)
plt.rcParams['figure.dpi'] = 120
plt.rcParams['font.size'] = 8
plt.rcParams['axes.titlesize'] = 9
plt.rcParams['axes.labelsize'] = 8
plt.rcParams['xtick.labelsize'] = 7
plt.rcParams['ytick.labelsize'] = 7
plt.rcParams['figure.autolayout'] = True
def intensity_transform(im, breakpoints):
    x = breakpoints[:, 0]
    y = breakpoints[:, 1]
    lut = np.interp(np.arange(256), x, y)
    lut = np.clip(lut, 0, 255).astype(np.uint8)
    transformed_im = lut[im]
    return transformed_im
eps = 1e-5
breakpoints = np.array([
    [0,0],
    [50,50],
    [50,100],
    [150,255],
    [150,150],
    [255,255],
])
im1 = cv2.imread('Images/emma.jpg', cv2.IMREAD_GRAYSCALE)
assert im1 is not None
intensity_transformed_im1 = intensity_transform(im1, breakpoints)
fig, axes = plt.subplots(1, 3, figsize=(14, 4.5), layout="constrained")
axes[0].imshow(im1, cmap="gray", vmin=0, vmax=255)
axes[0].set_title("Original Image")
axes[0].axis("off")
axes[1].imshow(intensity_transformed_im1, cmap="gray", vmin=0, vmax=255)
axes[1].set_title("Transformed Image (Set 1)")
axes[1].axis("off")
axes[2].plot(breakpoints[:, 0], breakpoints[:, 1], "ro-", linewidth=2, markersize=10)
axes[2].set_xlim(0, 255)
axes[2].set_ylim(0, 255)
axes[2].set_xlabel("Input Intensity")
axes[2].set_ylabel("Output Intensity")
axes[2].set_title("Intensity Transformation (Given Breakpoints)")
axes[2].grid(True, linestyle="--", alpha=0.6)
plt.tight_layout()
plt.show()
breakpoints_new = np.array([
```

```

[0, 0],
[60, 100],
[160, 200],
[255, 255]
])
intensity_transformed_im2 = intensity_transform(im1, breakpoints_new)
fig, axes = plt.subplots(1, 3, figsize=(14, 4.5))
axes[0].imshow(im1, cmap="gray", vmin=0, vmax=255)
axes[0].set_title("Original Image")
axes[0].axis("off")
axes[1].imshow(intensity_transformed_im2, cmap="gray", vmin=0, vmax=255)
axes[1].set_title("Transformed Image (Set 2)")
axes[1].axis("off")
axes[2].plot(breakpoints_new[:, 0], breakpoints_new[:, 1], "ro-", linewidth=2, m
axes[2].set_xlim(0, 255)
axes[2].set_ylim(0, 255)
axes[2].set_xlabel("Input Intensity")
axes[2].set_ylabel("Output Intensity")
axes[2].set_title("Intensity Transformation (My Breakpoints)")
axes[2].grid(True, linestyle="--", alpha=0.6)
plt.tight_layout()
plt.show()

```

C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\339701235.py:46: UserWarning: The figure layout has changed to tight
plt.tight_layout()



The given transform has caused the middle range of pixels to have higher intensities, while keeping the highest range of pixels unchanged. In my tuned transform, I have adjusted the intensity transform to brighten the middle range of pixels without having any jumps, to keep the image looking visually pleasing.

```

In [18]: im2 = cv2.imread('Images/brain.png', cv2.IMREAD_GRAYSCALE)
assert im2 is not None
breakpoints_white = np.array([




```

```

    [0, 0],
    [50, 50],
    [100, 50],
    [120, 255],
    [200, 50],
    [255, 50],
])
breakpoints_grey = np.array([
    [0, 0],
    [50, 50],
    [100, 50],
    [150, 50],
    [220, 255],
    [255, 50],
])
white_matter = intensity_transform(im2, breakpoints_white)
grey_matter = intensity_transform(im2, breakpoints_grey)
fig, axes = plt.subplots(2, 2, figsize=(10, 9), layout="constrained")
axes[0, 0].imshow(white_matter, cmap="gray", vmin=0, vmax=255)
axes[0, 0].set_title("White Matter")
axes[0, 0].axis("off")
axes[0, 1].plot(breakpoints_white[:, 0], breakpoints_white[:, 1], "ro-")
axes[0, 1].set_xlim(0, 255)
axes[0, 1].set_ylim(0, 255)
axes[0, 1].set_xlabel("Input Intensity")
axes[0, 1].set_ylabel("Output Intensity")
axes[0, 1].set_title("Intensity Transformation (White Matter)")
axes[0, 1].grid(True, linestyle="--", alpha=0.6)
axes[1, 0].imshow(grey_matter, cmap="gray", vmin=0, vmax=255)
axes[1, 0].set_title("Grey Matter")
axes[1, 0].axis("off")
axes[1, 1].plot(breakpoints_grey[:, 0], breakpoints_grey[:, 1], "ro-")
axes[1, 1].set_xlim(0, 255)
axes[1, 1].set_ylim(0, 255)
axes[1, 1].set_xlabel("Input Intensity")
axes[1, 1].set_ylabel("Output Intensity")
axes[1, 1].set_title("Intensity Transformation (Grey Matter)")
axes[1, 1].grid(True, linestyle="--", alpha=0.6)
plt.tight_layout()
plt.show()

```

C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\3069350537.py:42: UserWarning: The figure layout has changed to tight
 plt.tight_layout()



Here I adjusted the intensities to peak at the optimal intensity for highlighting the White and Grey matter, at 120 and 220 intensity values respectively. I found these values by manually adjusting the transform while checking for visual clarity of the required sections of the image. Evidently, the simple intensity transform has also amplified the noise presence in both the images, since noise is present in the entire range of intensities.


```
In [ ]: im3 = cv2.imread('Images/feet.jpeg')
assert im3 is not None
im3_rgb = cv2.cvtColor(im3, cv2.COLOR_BGR2RGB)
im3_lab = cv2.cvtColor(im3, cv2.COLOR_BGR2LAB)
l_channel_im3 = im3_lab[:, :, 0]
def gamma_correction(image, gamma):
    return np.clip(255 * (image / 255) ** gamma, 0, 255).astype(np.uint8)
l_channel_im3_corrected = gamma_correction(l_channel_im3, 0.6)
im3_lab_corrected = im3_lab.copy()
im3_lab_corrected[:, :, 0] = l_channel_im3_corrected
im3_corrected = cv2.cvtColor(im3_lab_corrected, cv2.COLOR_LAB2RGB)
channels = ("L (Lightness)", "a (Green-Red)", "b (Blue-Yellow)")
colors = ("black", "magenta", "orange")
fig, axes = plt.subplots(1, 4, figsize=(16, 4), layout="constrained")
axes[0].imshow(im3_rgb)
axes[0].set_title("Original Image")
axes[0].axis("off")
```

```

for idx, (name, color) in enumerate(zip(channels, colors)):
    ax = axes[idx + 1]
    hist = cv2.calcHist([im3_lab], [idx], None, [256], [0, 256])
    ax.plot(hist, color=color, lw=1.5)
    ax.set_title(f"{name} Channel")
    ax.set_xlim([0, 256])
    ax.set_yscale("log")
    ax.set_xlabel("Pixel Value (8-bit)")
    ax.grid(alpha=0.3)
axes[1].set_ylabel("Pixel Count")
plt.tight_layout()
plt.show()
fig, axes = plt.subplots(1, 4, figsize=(16, 4))
axes[0].imshow(im3_corrected)
axes[0].set_title("Corrected Image (Gamma = 0.6)")
axes[0].axis("off")
for idx, (name, color) in enumerate(zip(channels, colors)):
    ax = axes[idx + 1]
    hist = cv2.calcHist([im3_lab_corrected], [idx], None, [256], [0, 256])
    ax.plot(hist, color=color, lw=1.5)
    ax.set_title(f"{name} Channel")
    ax.set_xlim([0, 256])
    ax.set_yscale("log")
    ax.set_xlabel("Pixel Value (8-bit)")
    ax.grid(alpha=0.3)
axes[1].set_ylabel("Pixel Count")
plt.show()

```

C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\3123337396.py:28: UserWarning: The figure layout has changed to tight
plt.tight_layout()



The gamma correction has shifted and adjusted the histogram of the L channel to the right, which has brightened the image and made the white matter more visible. The gamma value of 0.6 was chosen after trial and error to find the most visually pleasing image. The gamma correction function has also amplified the noise already present in the L plane (visible in the original histogram), causing more significant noise spikes due to the non-linear nature of gamma correction function.

```

In [21]: im4 = cv2.imread("Images/spiderman.jpeg")
         assert im4 is not None

```


```



im4_rgb = cv2.cvtColor(im4, cv2.COLOR_BGR2RGB)
im4_hsv = cv2.cvtColor(im4, cv2.COLOR_BGR2HSV)
def f_transform(image, a, sigma=70):
    return np.clip(image + 128.0 * a * np.exp(-((image - 128.0) ** 2) / (2 * sigma**2)), 0, 255)
im4_saturation = im4_hsv[:, :, 1]
## Value of a was chosen as a = 0.6 from trial and error for visually pleasing result
im4_saturation_transform = f_transform(im4_saturation, a=0.6, sigma=70)
x = np.arange(256, dtype=np.float32)
a = 0.6
sigma = 70
fx = np.minimum(x + a * 128.0 * np.exp(-((x - 128.0) ** 2) / (2.0 * sigma**2)), 255)
im4_corrected = im4_hsv.copy()
im4_corrected[:, :, 1] = im4_saturation_transform
im4_corrected = cv2.cvtColor(im4_corrected, cv2.COLOR_HSV2RGB)
fig, axes = plt.subplots(1, 4, figsize=(16, 4), layout="constrained")
axes[0].imshow(im4_hsv[:, :, 0], cmap="hsv")
axes[0].set_title("Hue of Image")
axes[0].axis("off")
axes[1].imshow(im4_hsv[:, :, 1], cmap="hsv")
axes[1].set_title("Saturation of Image")
axes[1].axis("off")
axes[2].imshow(im4_hsv[:, :, 2], cmap="hsv")
axes[2].set_title("Value of Image")
axes[2].axis("off")
axes[3].imshow(im4_saturation_transform, cmap="hsv")
axes[3].set_title("Transformed Saturation Channel")
axes[3].axis("off")
plt.tight_layout()
plt.show()
fig, axes = plt.subplots(1, 3, figsize=(14, 4.5))
axes[0].plot(x, fx, color="crimson", lw=2, label=f"Transform ($a={a}$, $\sigma={sigma}$)")
axes[0].set_title("Intensity Transformation Function for Saturation", fontsize=12)
axes[0].set_xlabel("Input Intensity $x$", fontsize=10)
axes[0].set_ylabel("Output Intensity $f(x)$", fontsize=10)
axes[0].set_xlim([0, 255])
axes[0].set_ylim([0, 255])
axes[0].grid(alpha=0.3)
axes[0].legend(loc="upper left")
axes[0].set_aspect("equal", adjustable="box")
axes[1].imshow(im4_rgb)
axes[1].set_title("Original Image")
axes[1].axis("off")
axes[2].imshow(im4_corrected)
axes[2].set_title("Vibrance-enhanced Image")
axes[2].axis("off")
plt.tight_layout()
plt.show()

```

C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\1858499124.py:30: UserWarning: The figure layout has changed to tight

```
plt.tight_layout()
```






The given transform function was applied with value $a = 0.6$ applied onto the saturation channel of the HSV image. The value of a was chosen after trial and error to find the most visually pleasing image. The transform has brightened the saturation channel, which has made the colors in the image more vivid and vibrant (since the Saturation channel consists of the intense color regions of the image).

```
In [22]: im5 = cv2.imread("Images/shells.tif", cv2.IMREAD_GRAYSCALE)
assert im5 is not None
def hist_equal(image):
    L = 256
    M, N = image.shape
    hist = cv2.calcHist([image], [0], None, [L], [0, L])
    cdf = np.cumsum(hist)
    transform_map = np.array([(L-1)/(M*N)*cdf[k] for k in range (256)]).astype("float")
    equalized_image = transform_map[image]
    return equalized_image
im5_equalized = hist_equal(im5)
hist_orig = cv2.calcHist([im5], [0], None, [256], [0, 256]).ravel()
hist_eq = cv2.calcHist([im5_equalized], [0], None, [256], [0, 256]).ravel()
bins = np.arange(256)
fig, axes = plt.subplots(1, 4, figsize=(16, 4), layout="constrained")
axes[0].imshow(im5, cmap="gray", vmin=0, vmax=255)
axes[0].set_title("Original Image")
axes[0].axis("off")
axes[1].bar(bins, hist_orig, width=1.0, color="black", edgecolor="none")
axes[1].set_title("Histogram Before Equalization")
axes[1].set_xlim([0, 256])
axes[1].set_xlabel("Pixel Value (8-bit)")
axes[1].set_ylabel("Pixel Count")
axes[1].grid(alpha=0.3)
axes[2].imshow(im5_equalized, cmap="gray", vmin=0, vmax=255)
axes[2].set_title("Equalized Image")
axes[2].axis("off")
axes[3].bar(bins, hist_eq, width=1.0, color="crimson", edgecolor="none")
axes[3].set_title("Histogram After Equalization")
axes[3].set_xlim([0, 256])
axes[3].set_xlabel("Pixel Value (8-bit)")
axes[3].set_ylabel("Pixel Count")
axes[3].grid(alpha=0.3)
plt.tight_layout()
plt.show()
```

C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\2646575189.py:34: UserWarning: The figure layout has changed to tight
plt.tight_layout()




Histogram equalisation has enhanced the brightness of regions of the image by shifting a uniform distribution of the pixel intensities. This has also amplified the noise present in the image however, since noise is present in the entire range of intensities. There's also a loss of contrast in the image, since the pixel intensities have been shifted to a uniform distribution.

```
In [23]: im6 = cv2.imread("Images/woman.jpeg")
assert im6 is not None

im6_rgb = cv2.cvtColor(im6, cv2.COLOR_BGR2RGB)
im6_hsv = cv2.cvtColor(im6, cv2.COLOR_BGR2HSV)
h_plane, s_plane, v_plane = cv2.split(im6_hsv)
_, mask = cv2.threshold(s_plane, 15, 255, cv2.THRESH_BINARY)
mask_inv = cv2.bitwise_not(mask)
hist_foreground = cv2.calcHist([v_plane], [0], mask, [256], [0, 256])
cumsum = np.cumsum(hist_foreground)
foreground_size = cumsum[-1]
L = 256
transform_map = np.array(((L - 1) / foreground_size) * cumsum).astype(np.uint8)
v_equalized = transform_map[v_plane]
v_foreground_new = cv2.bitwise_and(v_equalized, v_equalized, mask=mask)
v_background = cv2.bitwise_and(v_plane, v_plane, mask=mask_inv)
final_v = cv2.add(v_background, v_foreground_new)
final_hsv = cv2.merge([h_plane, s_plane, final_v])
final_rgb = cv2.cvtColor(final_hsv, cv2.COLOR_HSV2RGB)
fig, axes = plt.subplots(2, 3, figsize=(14, 8), layout="constrained")
axes[0, 0].imshow(h_plane, cmap="gray")
axes[0, 0].set_title("Hue Plane")
axes[0, 0].axis("off")
axes[0, 1].imshow(s_plane, cmap="gray")
axes[0, 1].set_title("Saturation Plane")
axes[0, 1].axis("off")
axes[0, 2].imshow(v_plane, cmap="gray")
axes[0, 2].set_title("Value Plane")
axes[0, 2].axis("off")
axes[1, 0].imshow(mask, cmap="gray")
axes[1, 0].set_title("Foreground Mask")
axes[1, 0].axis("off")
axes[1, 1].imshow(im6_rgb)
axes[1, 1].set_title("Original Image")
axes[1, 1].axis("off")
axes[1, 2].imshow(final_rgb)
axes[1, 2].set_title("Result with Equalized Foreground")
axes[1, 2].axis("off")
plt.tight_layout()
plt.show()
```

```
C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\3284099365.py:39: UserWarning: The figure layout has changed to tight
plt.tight_layout()
```




I selected the saturation plane as the ideal one for foreground mask since it contains most of the intensity values needed from the foreground subject. After extracting the mask, I applied histogram equalisation on the Value plane, to ensure that only the brightness values are equalised across the foreground (instead of equalising the color intensities and distorting the image).

```
In [24]: im7 = cv2.imread("Images/einstein.jpeg", cv2.IMREAD_GRAYSCALE)
assert im7 is not None

sobel = np.array([[1,0,-1], [2, 0, -2], [1, 0, -1]], dtype=np.float32)
filtered1 = cv2.filter2D(im7, cv2.CV_32F, sobel)
[H, W] = np.shape(im7)
padded = np.pad(im7, pad_width=1, mode="constant", constant_values=0)
filtered = np.zeros((H, W), dtype=np.float32)
for i in range(H):
    for j in range(W):
        region = padded[i : i + 3, j : j + 3]
        filtered[i, j] = np.sum(region * sobel)
sobel2 = np.array([[1,0,-1]], dtype=np.float32)
sobel1 = np.array([[1], [2], [1]], dtype=np.float32)
filtered2 = cv2.filter2D(cv2.filter2D(im7, cv2.CV_32F, sobel1), cv2.CV_32F, sobel2)
fig, ax = plt.subplots(1, 3, figsize=(14, 4.5), layout="constrained")
ax[0].imshow(filtered1, cmap='gray')
ax[0].set_title('Sobel X Filter2D (a)')
ax[0].axis("off")
ax[1].imshow(filtered, cmap='gray')
ax[1].set_title('Sobel X Direct Method (b)')
ax[1].axis("off")
ax[2].imshow(filtered2, cmap='gray')
ax[2].set_title('Sobel X Separable Filter2D (c)')
ax[2].axis("off")
plt.tight_layout()
plt.show()
```

```
C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\2056982348.py:26: UserWarning: The figure layout has changed to tight
plt.tight_layout()
```



There is no visually noticeable difference between the SobelX filtering methods, but method (c), involving the decomposition of Sobel Horizontal into two vectors will have significantly reduced number of computations needed for the filtering output. Note that I padded the image when performing the direct filtering approach to avoid any errors or artifacts.

```
In [ ]: im8_1_big = cv2.imread("Images/im01.png")
im8_2_big = cv2.imread("Images/im02.png")
im8_3_big = cv2.imread("Images/im03.png")
im8_4_big = cv2.imread("Images/im04.jpg")x
im8_1_small = cv2.imread("Images/im01small.png")
im8_2_small = cv2.imread("Images/im02small.png")
im8_3_small = cv2.imread("Images/im03small.png")
im8_4_small = cv2.imread("Images/im04small.jpg")
def zoomNN(img, scale):
    return cv2.resize(img, None, fx=scale, fy=scale, interpolation=cv2.INTER
def zoomBilinear(img, scale):
    return cv2.resize(img, None, fx=scale, fy=scale, interpolation=cv2.INTER
def ssd(img1, img2):
    diff = img1.astype(np.float32) - img2.astype(np.float32)
    return np.mean(diff ** 2)

im8_1_zoomed_nn = zoomNN(im8_1_small, scale=4)
im8_1_zoomed_bilinear = zoomBilinear(im8_1_small, scale=4)
nn_SSD1 = ssd(im8_1_big, im8_1_zoomed_nn)
bilinear_SSD1 = ssd(im8_1_big, im8_1_zoomed_bilinear)
print('Normalized SSD for Nearest Neighbour of Image 1: ', nn_SSD1)
print('Normalized SSD for Bilinear of Image 1: ', bilinear_SSD1)

im8_2_zoomed_nn = zoomNN(im8_2_small, scale=4)
im8_2_zoomed_bilinear = zoomBilinear(im8_2_small, scale=4)
nn_SSD2 = ssd(im8_2_big, im8_2_zoomed_nn)
bilinear_SSD2 = ssd(im8_2_big, im8_2_zoomed_bilinear)
print('Normalized SSD for Nearest Neighbour of Image 2: ', nn_SSD2)
print('Normalized SSD for Bilinear of Image 2: ', bilinear_SSD2)

im8_3_zoomed_nn = zoomNN(im8_3_small, scale=4)
im8_3_zoomed_bilinear = zoomBilinear(im8_3_small, scale=4)
# %% Image 3 cannot be zoomed by the scaling factor 4 due to its size being indi
# nn_SSD3 = ssd(im8_3_big, im8_3_zoomed_nn)
# bilinear_SSD3 = ssd(im8_3_big, im8_3_zoomed_bilinear)
# print('Normalized SSD for Nearest Neighbour of Image 3: ', nn_SSD3)
# print('Normalized SSD for Bilinear of Image 3: ', bilinear_SSD3)

im8_4_zoomed_nn = zoomNN(im8_4_small, scale=4)
im8_4_zoomed_bilinear = zoomBilinear(im8_4_small, scale=4)
```

```

nn_SSD4 = ssd(im8_4_big, im8_4_zoomed_nn)
bilinear_SSD4 = ssd(im8_4_big, im8_4_zoomed_bilinear)
print('Normalized SSD for Nearest Neighbour of Image 4: ', nn_SSD4)
print('Normalized SSD for Bilinear of Image 4: ', bilinear_SSD4)

```

```

Normalized SSD for Nearest Neighbour of Image 1: 136.26904
Normalized SSD for Bilinear of Image 1: 115.0919
Normalized SSD for Nearest Neighbour of Image 2: 26.446089
Normalized SSD for Bilinear of Image 2: 18.345919
Normalized SSD for Nearest Neighbour of Image 4: 228.62233
Normalized SSD for Bilinear of Image 4: 197.4095

```

The zooming in was unsuccessful for the Third Image since the image is of a size 1459, which is not divisible by 4, causing the zoom by scale factor 4 to be impossible. In order to solve this, the only solution is to crop the image, or scale by a different factor.

```

In [16]: fig, axes = plt.subplots(3, 3, figsize=(14, 12), layout="constrained")

# Row 0: Image 1
axes[0, 0].imshow(cv2.cvtColor(im8_1_big, cv2.COLOR_BGR2RGB))
axes[0, 0].set_title("Image 1 Original", fontsize=11)
axes[0, 0].axis("off")

axes[0, 1].imshow(cv2.cvtColor(im8_1_zoomed_nn, cv2.COLOR_BGR2RGB))
axes[0, 1].set_title(f"NN Zoomed (x4)\nSSD: {nn_SSD1:.2f}", fontsize=11)
axes[0, 1].axis("off")

axes[0, 2].imshow(cv2.cvtColor(im8_1_zoomed_bilinear, cv2.COLOR_BGR2RGB))
axes[0, 2].set_title(f"Bilinear Zoomed (x4)\nSSD: {bilinear_SSD1:.2f}", fontsize=11)
axes[0, 2].axis("off")

# Row 1: Image 2
axes[1, 0].imshow(cv2.cvtColor(im8_2_big, cv2.COLOR_BGR2RGB))
axes[1, 0].set_title("Image 2 Original", fontsize=11)
axes[1, 0].axis("off")

axes[1, 1].imshow(cv2.cvtColor(im8_2_zoomed_nn, cv2.COLOR_BGR2RGB))
axes[1, 1].set_title(f"NN Zoomed (x4)\nSSD: {nn_SSD2:.2f}", fontsize=11)
axes[1, 1].axis("off")

axes[1, 2].imshow(cv2.cvtColor(im8_2_zoomed_bilinear, cv2.COLOR_BGR2RGB))
axes[1, 2].set_title(f"Bilinear Zoomed (x4)\nSSD: {bilinear_SSD2:.2f}", fontsize=11)
axes[1, 2].axis("off")


# Row 2: Image 4
axes[2, 0].imshow(cv2.cvtColor(im8_4_big, cv2.COLOR_BGR2RGB))
axes[2, 0].set_title("Image 4 Original", fontsize=11)
axes[2, 0].axis("off")

axes[2, 1].imshow(cv2.cvtColor(im8_4_zoomed_nn, cv2.COLOR_BGR2RGB))
axes[2, 1].set_title(f"NN Zoomed (x4)\nSSD: {nn_SSD4:.2f}", fontsize=11)
axes[2, 1].axis("off")

axes[2, 2].imshow(cv2.cvtColor(im8_4_zoomed_bilinear, cv2.COLOR_BGR2RGB))
axes[2, 2].set_title(f"Bilinear Zoomed (x4)\nSSD: {bilinear_SSD4:.2f}", fontsize=11)
axes[2, 2].axis("off")

plt.show()

```




```
In [25]: im9 = cv2.imread("Images/flower.jpeg")
assert im9 is not None

im9_rgb = cv2.cvtColor(im9, cv2.COLOR_BGR2RGB)
H, W = im9.shape[:2]
mask = np.zeros((H, W), np.uint8)
bgdModel = np.zeros((1, 65), np.float64)
fgdModel = np.zeros((1, 65), np.float64)
rect = (10, 150, 550, 500)
cv2.grabCut(im9, mask, rect, bgdModel, fgdModel, iterCount=5, mode=cv2.GC_INIT_W
fg_mask = np.where((mask == cv2.GC_FGD) | (mask == cv2.GC_PR_FGD), 255, 0).astype
bg_mask = cv2.bitwise_not(fg_mask)
foreground = cv2.bitwise_and(im9_rgb, im9_rgb, mask=fg_mask)
background = cv2.bitwise_and(im9_rgb, im9_rgb, mask=bg_mask)
blurred_bg = cv2.GaussianBlur(background, (51, 51), sigmaX=0)
enhanced_image = cv2.add(foreground, blurred_bg)

fig, axes = plt.subplots(1, 5, figsize=(18, 3.8), layout="constrained")
axes[0].imshow(im9_rgb)
axes[0].set_title("Original Image")
axes[0].axis("off")
axes[1].imshow(fg_mask, cmap="gray")
axes[1].set_title("Final Segmentation Mask")
axes[1].axis("off")
axes[2].imshow(foreground)
axes[2].set_title("Foreground Image")
axes[2].axis("off")
axes[3].imshow(background)
axes[3].set_title("Background Image")
axes[3].axis("off")
axes[4].imshow(enhanced_image)
```

```
axes[4].set_title("Enhanced Image (Blurred Background)")
axes[4].axis("off")
plt.tight_layout()
plt.show()
```

C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\498482814.py:34: UserWarning: The figure layout has changed to tight
plt.tight_layout()



For the GrabCut algorithm, I manually adjusted the region for the flower using the defined rect region (I wanted to select only the main flower and exclude the flower bud at the bottom of the image). When we blur the background, the region without the foreground is pure black pixels. These will get blurred with the surrounding regions, cause the dark edges present in the enhanced image.

In [26]:

```
im10 = cv2.imread("Images/mountain.jpeg")
assert im10 is not None
im10_rgb = cv2.cvtColor(im10, cv2.COLOR_BGR2RGB)
standard_gaussian_blur = cv2.GaussianBlur(im10_rgb, (51, 51), sigmaX=0)
standard_bilateral_filter = cv2.bilateralFilter(im10_rgb, d=7, sigmaColor=75, si
fig, axes = plt.subplots(1, 3, figsize=(14, 4.5), layout="constrained")
axes[0].imshow(im10_rgb, cmap="gray")
axes[0].set_title("Original Image")
axes[0].axis("off")
axes[1].imshow(standard_gaussian_blur)
axes[1].set_title("Gaussian Blur")
axes[1].axis("off")
axes[2].imshow(standard_bilateral_filter)
axes[2].set_title("Bilateral Filter")
axes[2].axis("off")
plt.tight_layout()
plt.show()
```

C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\2594712141.py:16: UserWarning: The figure layout has changed to tight
plt.tight_layout()



```

In [10]: def bilateral_filter_manual(image, radius, sigma_s, sigma_r):
    img = image.astype(np.float64)
    H, W = img.shape
    filtered = np.zeros((H, W), dtype=np.float64)
    for y in range(H):
        for x in range(W):
            I_x = img[y, x]
            weight_sum = 0.0
            pixel_acc = 0.0
            for dy in range(-radius, radius + 1):
                for dx in range(-radius, radius + 1):
                    yi = y + dy
                    xi = x + dx
                    if 0 <= yi < H and 0 <= xi < W:
                        I_xi = img[yi, xi]
                        spatial_dist_sq = dx * dx + dy * dy
                        g_s = math.exp(-spatial_dist_sq / (2.0 * sigma_s * sigma_s))
                        range_dist_sq = (I_xi - I_x) ** 2
                        g_r = math.exp(-range_dist_sq / (2.0 * sigma_r * sigma_r))
                        w = g_s * g_r
                        weight_sum += w
                        pixel_acc += w * I_xi
            filtered[y, x] = pixel_acc / weight_sum if weight_sum > 0 else I_x
    return np.clip(np.round(filtered), 0, 255).astype(np.uint8)

im10_r = im10_rgb[:, :, 0]
im10_r_bilateral = bilateral_filter_manual(im10_r, radius=7, sigma_s=75, sigma_r=75)
im10_g = im10_rgb[:, :, 1]
im10_g_bilateral = bilateral_filter_manual(im10_g, radius=7, sigma_s=75, sigma_r=75)
im10_b = im10_rgb[:, :, 2]
im10_b_bilateral = bilateral_filter_manual(im10_b, radius=7, sigma_s=75, sigma_r=75)

```

```

In [27]: merged = cv2.merge([im10_r_bilateral, im10_g_bilateral, im10_b_bilateral])
fig, axes = plt.subplots(1, 3, figsize=(14, 4.5), layout="constrained")
axes[0].imshow(im10_rgb, cmap="gray")
axes[0].set_title("Original Image")
axes[0].axis("off")
axes[1].imshow(standard_gaussian_blur)
axes[1].set_title("Gaussian Blur")
axes[1].axis("off")
axes[2].imshow(merged)
axes[2].set_title("Manually Calculated Bilateral Filter")
axes[2].axis("off")
plt.tight_layout()
plt.show()

```

C:\Users\Oshan\AppData\Local\Temp\ipykernel_7980\51781512.py:12: UserWarning: The figure layout has changed to tight
 plt.tight_layout()



```
In [15]: #I am using the SSD function from before to compare between OpenCV2 bilateral fi
         ssd_bilaterals = ssd(standard_bilateral_filter, merged)
         print("Normalized SSD between OpenCV2 Bilateral Filter and Manually Calculated B
```

Normalized SSD between OpenCV2 Bilateral Filter and Manually Calculated Bilateral Filter: 201.461

I decided to use the Sum of Squared Difference formula from the earlier questions to compare the efficacy of the Manual implementation and Direct implementation of Bilateral filtering (done on all three channels). Evidently, even with the same parameters, the manual calculation is not as good at filtering as the OpenCV2 function. I believe this is due to the fact that I am calculating the filter channelwise and combining, whereas the OpenCV2 function directly computes the joint distances on all channels together at once.